

# Parallel Architectures and Programming Models, 2023W

## **Assignment 1: OpenMP Tasking, Roofline Model**

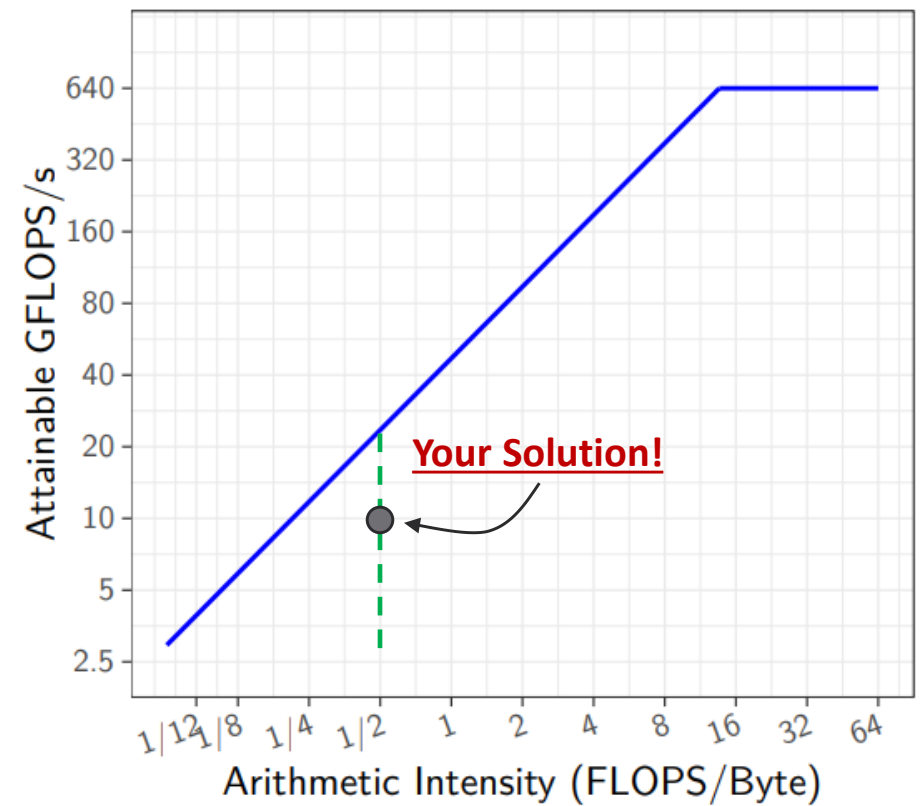
# GPSA :: ROOFLINE MODEL

Include the following in your report of assignment 1:

- 1) a **log-log** plot of the **Roofline Model** of a **single computational node of ALMA**;
- 2) include your **best performing** (32 cores) parallel implementation of GPSA to the model obtained in 1);

Based on the **Roofline Model obtained above**, and the **Arithmetic Intensity (AI)** of your solution, address the following questions in your report and **justify your answers**:

- 1) is your solution **memory-** or **compute-bounded**?
- 2) does it make sense to try to **optimize your parallel solution**? If yes, which **optimization techniques** could you employ?



# ABOUT ALMA SPECIFICATIONS...

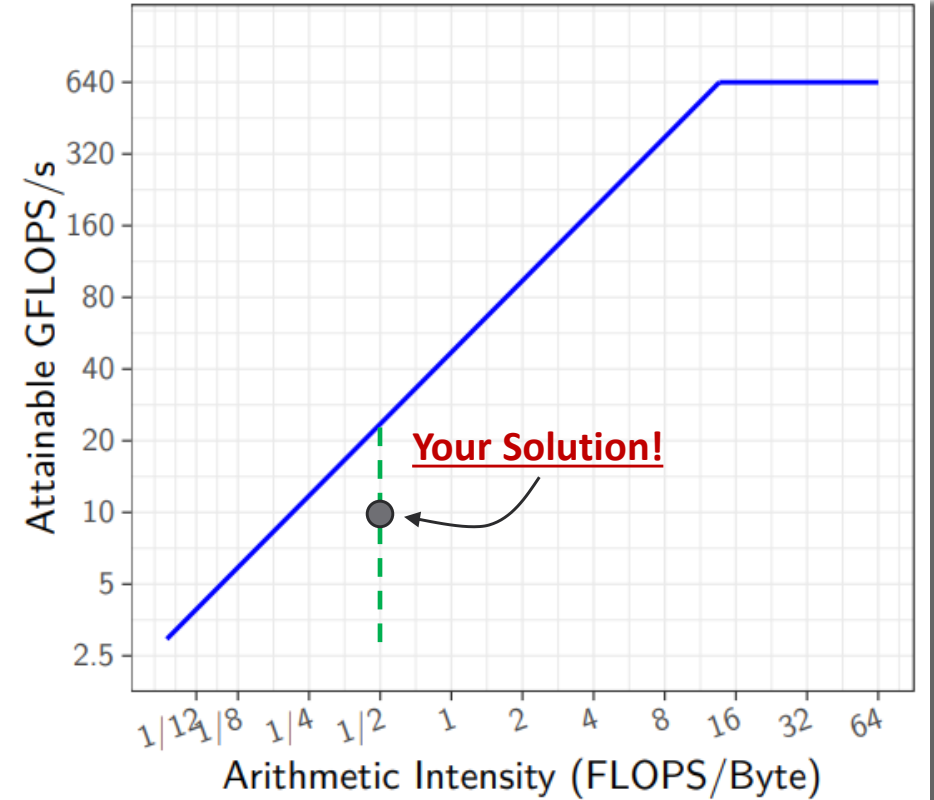
A **single computational node** of ALMA is composed by 2 Intel Xeon E5-2650 CPUs (with **2.0 GHz Clock Rate**).

Each CPU has **8 physical cores**, and each core can execute up to **2 AVX instructions per cycle**:

- `_mm_256_add_ps`: 8 single-precision FLOPs (ADD)
- `_mm_256_mul_ps`: 8 single-precision FLOPs (MUL)

Additionally, the STREAM benchmark shows that the **sustained memory bandwidth** of each computational node is equal to **57 GB/s**.

Note that the size of a single-precision floating-point number (`float`) is equal to **4 bytes**.



# ABOUT GPSA ARITHMETIC INTENSITY...

To compute the **Arithmetic Intensity (AI)** of your parallel solution (32 cores), you should only consider the code **inside the innermost-loop** of GPAS:

- for simplicity you should ignore initialization loops.
- **If you modified the innermost-loop, calculate the Arithmetic Intensity of the code on the left instead.**

You should also **only consider floating-point operations (FLOPs)**, and **ignore all** int and char operations such as:

- e.g., `X[i-1]` reads a char and `cmap.at` reads an int.

For simplicity, you should **ignore Read and Write operations** to the substitution matrix `SUB`.

Lastly, assume that `std::max({match, del, insert})` requires **2 FLOPs** per loop iteration.

```
for (unsigned int i = 1; i < rows; i++) {
    for (unsigned int j = 1; j < cols; j++) {
        //innermost-loop: BEGIN
        float match = S[i-1][j-1]
        + SUB[cmap.at(X[i-1])][cmap.at(Y[j-1])];
        float del = S[i - 1][j] + gap_penalty;
        float insert = S[i][j - 1] + gap_penalty;
        S[i][j] = std::max({match, del, insert});

        visited++;
        //innermost-loop: END
    }
}
```

# ADDITIONAL INFORMATION

- The Roofline part of the assignment accounts for 20% of Assignment 1 grade.
- All information required to plot the Roofline model is either provided:
  - in this document (“About ALMA specifications” ),
  - or can be obtained by running your parallel solution (#FLOPs executed, execution time, ...).
- You should NOT use profilers to obtain information (PAPI, PERF, LIKWID,...)!
- Include a table in your report with the values you used to plot the Roofline Model:
  - E.g., #FLOPs executed, execution time, PeakFLOPs/s, etc...
- The deadline for this part of the assignment is the same as the first part! (23 of November)